

Jiayi Chen

177 Yellowstone DR APT 308, Charlottesville, VA 22903, United States
jc4td@virginia.edu | (+1) 434-234-7342 | [Google Scholar](#) | [Linkedin](#)

Education

University of Virginia

Ph.D. in *Computer Science* (GPA 4.0/4.0)

United States

08/2018 – present

Xi'an Jiaotong University

M.S. in *Control Science and Engineering* (Specialization: *Artificial Intelligence and Robotics*)

China

09/2015 – 06/2018

Xi'an Jiaotong University

B.E. in *Automation Engineering*

China

09/2011 – 06/2015

Special Class for the Gifted Young of China (少年班)

09/2009 – 06/2011

Technical Skills

Programming Languages: Python (7+ years), C++, SQL, Java, Matlab

Tools & Platforms: PyTorch (5+ years), Tensorflow, Jax (2+ years), Flax, Haiku, Gym, Robotic Operating System

Deep Neural Networks: Large Language Models (e.g. BERT/ViT/T5/LayoutLM/GPT), Deep Generative Models (e.g., VAE/CGAN/Autoregressive/Diffusion), DNN/CNN/LSTM, Seq2Seq, Graph Neural Networks

Research Interests

Multimodal Machine Learning (Vision, Language, Audio, and Spatiotemporal data), Multimodal Large Language Models, Graph Mining, Federated Learning, Lifelong Learning, Model Compression (Pruning/Distillation)

Recent Publications (Since 2020)

- [1] **Jiayi Chen**, Kishlay Jha, Aidong Zhang. "Capacity-aware Lifelong Learning with Causal Knowledge Hypergraph Learning". *Under Review*, submitted September 2023.
- [2] **Jiayi Chen**, Zhe Zhao, Hussein Hazimeh, Shuo Yang, Aidong Zhang, Ed Chi. "Genes of Large Language Models: Toward Adversarial Knowledge Inheritance for Improved Distillation". *Manuscript*. 2023. *Work done during internship at Google
- [3] **Jiayi Chen**, Hanjun Dai, Bo Dai, Aidong Zhang, Wei Wei. "On Task-personalized Multimodal Few-shot Learning for Visually-rich Document Entity Retrieval". *Findings of Association for Computational Linguistics: EMNLP* 2023. (long paper)
*Work done during an internship at Google
- [4] Lijun Yu, Jin Miao, Xiaoyu Sun, **Jiayi Chen**, AG Hauptmann, Hanjun Dai, Wei Wei. "DocumentNet: Bridging the Data Gap in Document Pre-Training". The 2023 *Conference on EMNLP: Industry Track*. *Work done during internship at Google
- [3] **Jiayi Chen**, Aidong Zhang. "On hierarchical disentanglement of interactive behaviors for multimodal spatiotemporal data with incompleteness". The 29th *ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)*, 2023. (research track)
- [6] **Jiayi Chen**, Aidong Zhang. "FedMSplit: correlation-adaptive federated multitask learning across multimodal split networks". The 28th *ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)*, 2022. (research track)
- [7] **Jiayi Chen**, Aidong Zhang. "Topological transduction for hybrid few-shot learning". In *Proceedings of the ACM Web Conference (WWW)*, 2022. (research track)
- [8] **Jiayi Chen**, Aidong Zhang. "HetMAML: task-heterogeneous model-agnostic meta-learning for few-shot learning across modalities". *ACM International Conference on Information and Knowledge Management (CIKM)*, 2021. (full paper)
- [9] **Jiayi Chen**, Aidong Zhang. "HGMF: heterogeneous graph-based fusion for multimodal data with incompleteness". The 26th *ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)*, 2020. (research track)
- [10] Jiachen Yu, Li Jin, **Jiayi Chen**, Youzi Xiao, Zhiqiang Tian, and Xuguang Lan. "Deep semantic space guided multi-scale neural style transfer." *Multimedia Tools and Applications*, 2022, 81(3): 3915-3938. ([Journal](#))

Work Experience

Google DeepMind

Research Intern

Mountain View, CA

June 2023 – Sep 2023

- **Mentors:** Zhe Zhao ([Google DeepMind](#)), Shuo Yang ([YouTube](#)), Hussein Hazimeh ([Google Research](#))
- **Project:** [Topology-agnostic Knowledge Inheritance for Improved Knowledge Distillation](#). The task is to explore both *direct* and *indirect* knowledge from pre-trained Large Language Models (LLMs) for *recommendation* tasks and vision tasks. The objective is to harness Parameter-Level knowledge of LLMs (e.g., *pre-trained weights* and *their topological structure*) to *accelerate/enhance* the knowledge transfer at the Inference Level (often referred to as Distillation).
- **Accomplishment:** Successfully achieved (1) a **first-author** research paper (in review); (2) potential product improvement.

Google LLC

SWE Intern

Mountain View, CA

June 2022 – Dec 2022

- **Mentors:** Hanjun Dai ([Google DeepMind](#)), Bo Dai ([Google DeepMind](#)), Wei Wei ([Cloud DocAI & Core ML App](#))
- **Project:** [Few-shot Multimodal Entity Retrieval for Visually-rich Documents](#). Visually-rich documents consist of three modalities (*language*, *image*, and *layout structure* of contents). The task is to harness meta-knowledge to *accelerate* the learning process of (1) understanding *new* document types given a pre-trained **text-image Large Language Model**; (2) localizing *rarely-occurred* key information types from *out-of-distribution* information.
- **Accomplishment:** Successfully achieved **two** research papers: (1) published a **first-author** research paper in EMNLP 2023 Findings; and (2) contributed to experimental work for another research paper ([PDF](#)).

Selected Project Experience

On Multi-modality Understanding for Semi-/Unsupervised, Few-shot, or Generative Tasks

05/2019 – Present

- Proposed a graph-factorized **Gaussian Process VAE** to **disentangle** the underlying dynamic systems of partially observed **spatiotemporal interactions** between multiple subjects; and then robustly **generate** imputed data or future data based on learned disentangled representations.
- Proposed a heterogeneous hyper-node **graph** for the **multi-task** multimodal alignment task.
- Proposed a dynamic graph-based **federated learning** approach for **privacy-preserving personalized task adaptation** with heterogeneous domain shifts and modality types.

Toward Emotion Controllable Chatbot via Conditional Generative Models

11/2018 – 03/2019

- **Keywords:** NLP, Natural Language Generation, Reinforcement Learning, GANs, VAEs, Sequential Data Generation
- Generate texts/language data whose semantics or attributes (e.g., emotional tones) can be controlled; **adversarially** trained the sequential data generator with Conditional GAN and VAE, while using **reinforcement learning (RL)** to promote the authenticity of the generated language.

Knowledge-guided Neural Style Transfer of 3D Scene Models

03/2017 – 02/2018

- **Keywords:** Computer Vision, Neural Style Transfer
- Studied human knowledge-guided **neural style transfer**, focusing on improving the illusion of space in generated images by simulating how artists harness skills to understand and reproduce a 3D scene (e.g., geometric structures, lighting and shallow); also studied 3D non-photorealistic rendering based on the neural style transfer paradigm.
- Proposed an illumination-guided deep alignment method using CNN, Lighting Path Expression, and PatchMatch.
- Created a 2D-3D dataset, including 3D models rendered by multiple types of lighting (by Maya), 2D photos annotated by lighting and segmentation (by Photoshop and Matlab), and a hand-drawn stylistic material for testing (CorelPainter).

Vision-guided Drone Autonomous Navigation & Search-and-rescue System

01/2016 – 12/2016

- **Keywords:** Robot Manipulation, Robot Perception, 3D Computer Vision, C++, Robotic Operating System
- Developed a real-time **Vision**-based module (including **Object detection** and **3D localization**) on Drone's autonomous navigation system. Real-time 3D poses of objects are estimated from 2D images using the real-time camera gimbal data based on Homography and 3D Geometry.
- Developed an autonomous **tracking-and-landing** module for Drone landing on a moving vehicle with **safety guarantee**, and designed a platform to manage multimodal sensors and commands (*tools:* **Python**, Linux, ROS).
- Won the **4th** place from 130+ international teams in "2016 DJI Developer Challenge", NY, USA (as a team of three).

Selected Awards

- 2022 KDD Student Travel Award
- 2018 Outstanding Master's Thesis Award
- 2016 Industrial Scholarship (*Top 3%, Rank: 4/104*)
- **4th** place in world **Top10** in “2016 **International** Contest of DJI Drone Developer Challenge” (*Top 3%, Rank: 4/130*)
- **2nd** prize in “2016 China **National** Graduate Contest on Smart-City Technology and Creative Design”
- **1st** prize in “2015 China **National** Contest on Traffic-Scene Image Processing”

Relevant Courses

Machine Learning and Computer Vision (A+), Natural Language Processing (A+), Reinforcement Learning (A+), Database Systems (A), Graph Mining (A+), Statistical Learning and Graphical Models (A), Computer Graphics (A+), Software Artifacts (A), Autonomous Mobile Robot (A+), Data Structures and Algorithms (A+), Computer Network (A), etc.

Additional

Teaching Assistant: *Program and Data Representation* (Fall 2019, Spring 2020); *Digital Image Processing* (Fall 2015)

Peer Review Services: Reviewer/external reviewer for NeurIPS, ICLR, AAAI, IJCAI, CIKM, WSDM, SDM, BIBM, etc.

Startup Experience: Joined a university **Startup** team (MachInsight) from 2015 to 2016, focusing on dealing with **(1) Computer Vision**-related projects for industry; **(2) Drone** (a type of Autonomous Aerial Robot) related software product driven by Computer Vision; and **(3) Augmented Reality** related product. E.g., at 2015 XJTU University Anniversary, we developed an augmented reality iOS app using Java, Swift, and Unity3D; in 2016, we developed a communicative module for Drone for connecting with Ground Vehicle to achieve enhanced visual tracking and safe landing.